


Foundry Agent Service

How to build AI-driven applications that scale AI agents



Day 1 of 2 — Vision, Architecture & Live Demos

Delivered for RAI Pubblicità · Mauro Minella, Senior Cloud Solution Architect

AGENDA

Four themes, two breaks, one story

1

Vision & Architecture

Why agents now, Microsoft Foundry, and the map of where everything fits

 *break after*

2

Foundry Agent Service — Build

The managed runtime and building blocks for a single production agent

 *break after*

3

The four IQs — Ground

Work / Foundry / Web / Fabric IQ: enterprise knowledge for agents

4

Agent 365 — Govern

The control plane to publish, secure and manage a fleet of agents

Two half-days, two audiences, one journey

Day 1 — today (4h)

For everyone: leadership + technical teams. Vision, architecture and the opportunity, with short live demos after each concept. Goal: shared language and a clear map.

Day 2 — tomorrow (4h)

For technical roles: CSA, developers, data scientists. Deeper demos and four hands-on labs you can follow step by step on your own machine.

Deep-dived in separate dedicated sessions (named today, not exhausted):

Evaluation

Quality, groundedness, safety scoring

Content Understanding

Multimodal document & media parsing

Observability & Tracing

End-to-end agent telemetry

Hosted Agents / MCP / A2A

Runtime & protocol deep-dive

RHYTHM

Overview → demo → takeaway, every time

1

Overview

The concept in plain language, then the real technical detail behind it.



2

Live demo

Short and ready-to-deliver on Foundry / Azure — no data loading, no setup.



3

Takeaway

Why it matters for you — as a leader, or as a builder.

PROMISE

Engaging intro, then real technical depth. You will not leave with a marketing deck — you will leave able to reason about, and start building, production agents.

POSITIONING

The 2025 vision, in 2026's language

The program hasn't changed — the industry named it properly. Microsoft renamed the platform and the service; your three original pillars map cleanly onto today's architecture.

2025 proposal (original pillars)	2026 — current Microsoft naming
AI Foundry — Vision & Architecture	Microsoft Foundry (platform renamed from “Azure AI Foundry”)
AI Services in AI Foundry	Foundry Tools & the four IQs — cognitive services now consumed BY agents
AI Agent Service	Foundry Agent Service + Agent 365 (the fleet control plane)

HOW TO SAY IT TO THE ROOM


“We didn't change the plan — we're giving you 2026's correct terminology for the 2025 vision. Same destination, current map.”

1

PART 1

Vision & Architecture

Why agents now, *what* Microsoft Foundry is, and *the single map* that holds the whole day together.



THE SHIFT

From answering questions to completing work

- **A model answers. An agent acts — it plans, calls tools, checks results and iterates toward a goal.**
- The bottleneck moved from “model quality” to “how safely we let models take action on enterprise systems.”
- Reasoning models + open tool protocols (MCP, A2A) made reliable, multi-step automation practical in 2025–2026.
- The enterprise question is no longer “can it answer?” but “how do we run, ground and govern this at scale?”

1.3 billion

AI agents in enterprises projected by 2028

Source: IDC

TAKEAWAY

Agents are an operating-model change, not a feature. Today is about the platform that makes them safe and scalable.

EVOLUTION

Chatbot → Copilot → Agent → Agent fleet


Chatbot

Scripted Q&A. No memory, no tools.



Copilot

LLM assists a human in the loop.



Agent

Plans & acts on tools toward a goal.



Agent fleet

Many agents, governed like a workforce.



More autonomy · more value · more need for identity, grounding and governance →

What actually is an agent

Model

The reasoning engine, from the Foundry model catalog.

Instructions

System prompt / policy: role, goals, guardrails.

Tools

Functions, OpenAPI, MCP servers, A2A peers, code interpreter.

Knowledge

Grounding sources — files, search, the four IQs.

Memory

Short-term, long-term, procedural — managed for you.

Identity

Entra Agent ID — a first-class, governable identity.

The agentic loop: perceive → plan → call a tool → observe result → repeat until the goal is met

Microsoft Foundry — one place to build with AI

- **Model catalog** — frontier and open models, one API surface, deploy & compare.
- **Foundry Tools** — Speech, Vision, Content Safety, Content Understanding, Search.
- **Foundry Agent Service** — build, run and scale agents (our focus today).
- **Grounding & IQ** — connect enterprise knowledge to agents.
- **Evaluation, tracing & guardrails** — quality and safety built in.

One project, one identity model, one bill

- Portal + SDK + REST (v1)
- Azure networking, RBAC, Key Vault
- Entra identity end-to-end
- Region & data residency controls

The cognitive services didn't disappear — agents use them

Speech

STT / TTS, and Voice Live for real-time voice agents.

Vision

Image understanding, OCR, spatial analysis.

Content Safety

Moderation & abuse detection as a guardrail.


Content Understanding

Multimodal doc/media extraction (own session).

From “services you call” to “capabilities your agent has”

- In 2025 you called these services directly from application code.
- In 2026 an agent invokes them as tools — Speech becomes the agent's voice, Vision its eyes, Content Safety its guardrail.
- Same proven services; a new consumption pattern. This is exactly the old “AI Services” pillar, matured.

Build → Ground → Govern



One agent → a grounded agent → a governed fleet. Every topic today lives on exactly one of these planes.

Where everything fits — one reference map

GOVERN — Agent 365

Registry · Access control · Observability · Interoperability · Security · Publish to Teams & M365 Copilot

BUILD — Foundry Agent Service

Hosted & Prompt agents · Toolbox (MCP) · Memory · Voice Live · Entra Agent ID · Guardrails

GROUND — Knowledge & the four IQs

Work IQ · Foundry IQ · Web IQ · Fabric IQ + Foundry model catalog & Foundry Tools

Identity, security & telemetry run vertically through all three planes (Microsoft Entra · Defender · Purview)

An enterprise agent, end to end

- 1 User / App / Teams
- 2 Foundry Agent Service (hosted runtime)
- 3 Tools via Toolbox (MCP · OpenAPI · A2A)
- 4 Grounding: Foundry IQ / Fabric IQ
- 5 Entra Agent ID + Guardrails
- 6 Traced & governed in Agent 365

NEXT


This is exactly what we'll build up, plane by plane. First plane: BUILD — Foundry Agent Service. Let's take a short break, then dive in.

2

PART 2

Foundry Agent Service

The managed runtime and the building blocks for a single production-grade agent. This is the “Build” plane.



From “AI Agent Service” to “Foundry Agent Service”

- **Same service, renamed by Microsoft as the platform became “Microsoft Foundry”.**
- It grew from “an API to create an assistant” into a managed runtime for production agents.
- New since 2025:
 - hosted agents,
 - the Toolbox MCP endpoint,
 - built-in memory,
 - Entra Agent ID,
 - Voice Live.
- Docs & APIs now live under learn.microsoft.com/azure/foundry/agents — with a versioned v1 REST surface.

WHY IT MATTERS

The rename signals maturity: from an experimentation API to an enterprise runtime you can standardise on.

Two ways to build:

Prompt agents

Hosted agents

(next slide)

Prompt-based vs Hosted agents

	Prompt-based agent	Hosted agent
Definition	Model + instructions + tools, declared in the service	Your own code / framework, deployed to a managed runtime
Best when	Standard orchestration, fast start	Custom logic, webhooks, non-OpenAI protocols
Frameworks	Built-in orchestration	Agent Framework, LangGraph, Semantic Kernel, custom
Runtime	Managed by Foundry	Session-isolated, autoscaling, zero idle cost
Deploy	Configure in portal / YAML / SDK	Deploy with one command

Rule of thumb: start prompt-based; move to hosted when you need to bring your own code.

Hosted agents — a runtime you don't have to run

Session isolation

Each conversation runs isolated — no cross-tenant / cross-session leakage.

Autoscaling

Scales with load automatically; you don't manage servers or containers.

Zero idle cost

No traffic, no compute charge — pay for work, not for waiting.

One-command deploy

Ship code to the managed runtime; rollbacks & versioning handled.

Managed cross-cutting

Containerisation, web server, security, instrumentation — all built in.

Toolbox access

Reaches Code Interpreter, Web Search, AI Search, OpenAPI, MCP, A2A.

It removes the plumbing that normally sits between a working prototype and a production agent.

Create a prompt agent in the Foundry portal

- In the Foundry portal, open a project → Agents → New agent.
- Pick a model, give it a name and clear instructions (role + guardrails).
- Add a built-in tool and test the agent live in the playground.
- Result: a working agent in ~3 minutes, no infrastructure.

ON SCREEN

From empty project to a responding, tool-using agent — entirely in the browser.

MCP · A2A · OpenAPI — no protocol lock-in

MCP

Model Context Protocol — the open standard for exposing tools & context to any LLM client.

A2A

Agent-to-Agent — agents discover and call other agents, across vendors.

OpenAPI


Bring any existing REST API as a tool, straight from its spec.

- Open protocols mean your tools and integrations aren't tied to one vendor's runtime.
- Reuse what you already have (REST APIs, internal MCP servers) instead of rebuilding it.
- Interoperability by design — a strict requirement for enterprise and public-sector buyers.

TOOLBOX

One MCP endpoint for all your tools

- **Curate reusable tools and expose them through a single MCP-compatible endpoint.**
- **Centralised authentication** — connections and secrets configured once, reused everywhere.
- **Versioned updates** — change a tool centrally; every agent picks up the new version.
- **Built-in Foundry tools** plus your **custom MCP connections** behind the same door.



Add a tool and see the Toolbox endpoint

- Open the agent from Demo 1 → Tools → add a connection.
- Add a built-in tool (e.g. Web Search) and/or a custom MCP server URL.
- Ask a question that forces a tool call; watch the tool invocation & result.
- Point out the single Toolbox MCP endpoint the agent talks to.



ON SCREEN

One agent, one endpoint, many tools — added without touching code.

The 'Create skill' dialog box contains the following fields:

- Name ***: extract-keywords
- Description ***: Extracts main keywords from a provided text
- Instruction ***: Extract the top 5 keywords from the given text

At the bottom are 'Create' and 'Cancel' buttons.



Built-in memory — no external database

Short-term

The working context of the current conversation / thread — carried automatically.

Long-term

Facts and preferences that persist across sessions, per user or per agent.

Procedural

Learned how-to: reusable steps and workflows the agent should follow.

- Managed by the service — you don't provision or operate a vector store or state DB.
- Threads persist conversation state; memory makes agents feel continuous, not amnesiac.
- Enterprise controls: scoping, retention and deletion respect your data governance.

Memory across turns

- In the playground, tell the agent a preference ("answer me in Italian, briefly").
- Continue the thread — the agent keeps the context without repeating it.
- Start a new thread; show what persists (long-term) vs what resets (short-term).
- No database was set up — memory is part of the service.

ON SCREEN

The agent remembers within a thread and recalls preferences across threads — zero infra.

```
# Procedural memory is enabled to allow the agent
# to remember and execute a series of steps for a specific task.

agent = Agent(
    model="gpt-4.1",
    enable_procedural_memory=True
)

agent.learn_procedure(
    name="analisi_campagna",
    steps=[
        "Cerca dati storici",
        "Genera un grafico",
        "Estrai keyword",
        "Scrivi un riassunto"
    ]
)
```

Procedural memory (sample)

Entra Agent ID — every agent is a first-class identity

- **Foundry auto-provisions an Entra identity for the agent across its lifecycle.**
- Three parts: blueprint (the app template) → instance (a deployment) → agent user (runtime identity in your org).
- Least-privilege RBAC & Conditional Access — the agent only reaches what it's granted.
- This is the bridge to the “Govern” plane: the same identity Agent 365 manages.

FOR GOVERNANCE & SECURITY

Agents stop being anonymous scripts. They get an auditable identity you manage like an employee's.

Blueprint → Instance → Agent user

Voice Live — real-time, speech-to-speech agents

- **Natural, low-latency speech-to-speech: the agent listens and talks in real time.**
- Powered by Azure Speech inside Foundry Tools; connect by agent ID — logic stays in the agent.
- The agent's prompt, tools and memory are reused; only the interaction channel changes.
- Ideal for contact-centre, accessibility and hands-free enterprise scenarios.

WHY IT LANDS

Same agent, now a voice. It makes “an agent” tangible for a non-technical audience in seconds.

Talk to your agent — Voice Live playground

- Open Foundry Tools → Speech → Voice Live playground; select the agent by ID.
- Speak a question out loud; the agent answers back in natural speech.
- Interrupt mid-sentence to show real-time turn-taking and low latency.
- Same agent, same tools — now voice-first, no extra build.

ON SCREEN

A live spoken conversation with the agent we built — straight from the portal playground.

Observability and guardrails — a first look

Built-in evaluators

Score groundedness, relevance, safety — automatically.

Trace linkage

Every run traced: prompt, tool calls, outputs, decisions.

Guardrails

Content Safety + policy enforcement on inputs & outputs.

Policy control

Consistent rules applied across agents.

SCOPE NOTE

Two dedicated sessions go deep on Evaluation and on Observability & Tracing. Today: know it exists and that it's built in — quality and safety are not bolt-ons.

Deep-dive: separate Evaluation session · separate Observability & Tracing session

Private networking, then a quick recap

Private networking (brief)

- Run agents inside your own network with bring-your-own VNet (BYO VNet).
- Minimise public egress; keep traffic on the private backbone.
- Combine with Private Endpoints, RBAC and Key Vault.
- Detail belongs in an architecture/networking deep-dive — flagged, not covered.

Build — recap

- Prompt & hosted agents on a managed, zero-idle runtime.
- Open standards (MCP/A2A/OpenAPI) via one Toolbox endpoint.
- Built-in memory; Voice Live; Entra Agent ID.
- Quality & safety built in — deep-dived separately.

 **Break 2 — then the “Ground” plane: the four IQs**

3

PART 3

The four IQs

Grounding agents in enterprise knowledge with Microsoft IQ: Work, Foundry, Web and Fabric IQ. This is the “Ground” plane.



A smart model is not an informed model

Without grounding

- Answers from training data only
- Confident but possibly wrong
- No access to your policies or data
- Hard to audit a claim

With grounding

- Retrieves your authoritative sources
- Answers with citations
- Reflects current, internal reality
- Every claim is traceable

The enterprise intelligence layer

Four complementary sources of context that together give agents a complete view of your organisation:

Work IQ

How your people work — the flow of work across Microsoft 365.

Foundry IQ

Your policies & authoritative documents — the knowledge agents cite.

Web IQ

Fresh context from the public web, safely.

Fabric IQ

Business entities & data — the semantic state of your business.

Not four products to stitch together — one intelligence layer agents draw on.

Context on how your people work

- **Grounds agents in the flow of work across Microsoft 365 — mail, chats, meetings, docs.**
- Lets an agent understand roles, relationships and the way work actually happens.
- Powers the interoperability of Agent 365 — agents that fit into human workflows.
- Respects existing Microsoft 365 permissions — agents see only what the user may see.

IN ONE LINE

Work IQ is why an agent can act like a well-onboarded colleague, not an outsider.

Your policies and authoritative documents

- **A knowledge base (built on Azure AI Search) agents query for grounded, cited answers.**
- Agentic retrieval: the agent plans queries, retrieves, and returns citation-backed responses.
- Connected to Foundry Agent Service via a project connection / MCP tool.
- Generally available in the 2026-04-01 REST API; newer capabilities in preview.

THIS IS THE DEMO

Foundry IQ is what we'll connect live: give the agent a knowledge source, then watch it cite.

Fresh context from the public web

- Gives agents safe, current information from the public web when internal data isn't enough.
- Useful for market, news, and reference lookups beyond the enterprise boundary.
- **Governed: results flow through the same safety and policy controls as any tool.**
- Complements — never replaces — your authoritative internal knowledge.

BALANCE
Web IQ for “what's happening out there”,
Foundry & Fabric IQ for “what's true in here”.

WEB IQ	Grounding for Bing Search
Citations are returned	Pages are returned
Optimized for Reasoning	Optimized for ranking
Interpretation	Raw HTML pages
Computer-opmized index	Human-opmized index
Ready-made answers	Needs LLM to interpret pages
Parallel execution (quicker)	Sequential analysis

The semantic state of your business

- **Grounds agents in business entities and data — customers, campaigns, revenue — not raw tables.**
- Elevates data into the language of your business via ontologies / a semantic layer.
- Built on Microsoft Fabric: analytical, real-time and operational data unified.
- Lets agents reason in business concepts, so answers align with your KPIs.

An agent that knows a “campaign” or “advertiser” as your business defines it — not just a column name.

A “campaign” is a business entity, not a table

THE BUSINESS ASKS

Which campaign had the best ROAS last quarter, counting only active customers?

WITHOUT FABRIC IQ

The agent can only stop and ask back

- There’s no ROAS column.
- No “active customer” field exists.
- Which join links Campaigns and Revenue?
- How is “active” defined?

WITH FABRIC IQ

The agent answers in your business terms

- **BackToSchool_2024 — highest ROAS at 3.4**
- €128,000 attributed revenue from active customers
- on €37,500 spend · +18% vs prior quarter
- Best-performing channel: Paid Search

WHY IT WORKS

Fabric IQ already knows what a campaign, an active customer (≥ 2 purchases in 90 days), an attributed revenue (within 7 days of a click) and ROAS are — and which tables compose them. The agent doesn’t just read data; it understands the business.

Ground the agent with Foundry IQ

- Add a Foundry IQ knowledge base as a project connection to the agent.
- Ask a question answerable only from that source.
- Show the answer with inline citations back to the documents.
- Ask something outside the source — the agent declines instead of guessing.

ON SCREEN


The same agent, now answering from your knowledge — with citations, not guesses.

4

PART 4

Agent 365

The control plane to publish, secure and manage a fleet of agents across the organisation. This is the “Govern” plane.



THE PROBLEM

Agent sprawl and shadow AI, at scale

- **If there really are ~1.3B agents by 2028, most enterprises will run hundreds internally.**
- Ungoverned agents = shadow AI: untracked identities, unclear ownership, hidden data access.
- IT systems built for humans aren't designed for autonomous actors that act at machine speed.
- The risk isn't one bad agent — it's having no map of what your agents can do.

**Manage agents
the way you
manage people**

THE SHIFT

Same trusted infrastructure — identity, security, admin — extended from people to agents.

The control plane for AI agents

1. **Governs agents from any source — Microsoft, open-source, or third-party platforms.**
2. Built on the trusted stack: Entra, Defender, Purview, Intune, the M365 admin center.
3. Generally available (from May 2026); licensed per user, part of the Microsoft 365 E7 suite.
4. Foundry agents integrate automatically — no separate registration step.

FOR LEADERSHIP

One place to see, secure and measure every agent — regardless of who built it or how.

CAPABILITIES

Five capabilities that make fleet-scale AI possible

1 · Registry

A single inventory of every agent — including shadow ones.

2 · Access control

Unique agent ID, least privilege, policy templates.

3 · Observability

Dashboards, telemetry, ROI and performance.

4 · Interoperability

Agents reach M365 apps & data via Work IQ.

5 · Security

Threat detection & remediation with Defender/Purview.


Observe · Govern · Secure — the three verbs of Agent 365.

One source of truth for every agent

- **A tenant-wide inventory: agents with Entra Agent ID, registered agents, and shadow agents.**
- Surfaced in the Microsoft 365 admin center; ownership, source and risk visible at a glance.
- Quarantine unsanctioned agents so they can't be discovered or reach resources.
- The Agent Store lets users find approved agents for their role in Teams & M365 Copilot.

KILLS THE BLIND SPOT

You can't govern what you can't see. The registry makes the whole fleet visible first.



Least privilege for autonomous actors

- Every agent has a unique agent ID — the identity we saw in the Build plane.
- Policy templates enforce a security baseline from day one.
- Entra applies adaptive, risk-based access; compromised agents are blocked automatically.
- Least privilege: an agent reaches only the resources its task requires.

BRIDGE TO BUILD

Remember Entra Agent ID? This is where that identity is governed at fleet scale.

 asbagent03

Available

 · Platform: Microsoft Foundry ·

+ Add tag

Details

 · Users · Data & tools · Security · Permissions · Activity

About this agent · Version 1.0

No description provided

Overview

Created	September 5, 2026	Last updated	September 5, 2026
Last published	—	Last used	Never
Channel	—	Publisher	None
Publisher type	Your users	Owner	Mauro Minella
Created by	Mauro Minella	Entra agent ID	a8cefb1-fe23-4a42-9947-b4f1...

Identity

Entra agent ID

a8cefb1-fe23-4a42-9947-b4f1640459e3

Blueprint ID

846b1837-7381-4e1a-94f4-5661792f3b9a

Source agent ID

/subscriptions/eca2eddb-0f0c-4351-a634-52751499eeea/providers/Microsoft.CognitiveServices/accounts/project/agents/asbagent03

Bot ID

—

Properties

Editors

1 users

Agent location

swedencentral

Model name

gpt-5.4-mini

Agentic type

Prompt agent

Agent Url

<https://foundry71...>

Foundry Azure resource ID



/subscriptions/eca2eddb-0f0c-4351-a634-52751499eeea/resourceGroups/rg-aifoundry7159/providers/Microsoft.CognitiveServices/accounts/project/agents/asbagent03

From monitoring to actionable insight

- **A unified dashboard maps connections among agents, users and resources.**
- Role-based reporting: IT, security and business leaders each see what matters to them.
- Performance & ROI: track speed, quality and business impact to guide investment.
- Compliance: logging, e-discovery and retention keep you audit-ready.

FOR LEADERS

Not just “are agents running?” but “are they worth it?” — measurable impact per agent.



Interoperability and Security

4 · Interoperability

- Agents use the same apps as people: Word, Excel, Outlook, SharePoint, Dynamics 365
- Work IQ gives agents work context to onboard into processes
- Build with Copilot Studio, Foundry, OSS or the Agent 365 SDK

5 · Security

- Defender detects threats & attacks that target agents
- Purview extends data protection & DLP to agent actions
- Entra blocks compromised agents; least-privilege enforced

One-click publishing to Teams & M365 Copilot

- **Publish an agent to Microsoft Teams and Microsoft 365 Copilot in one click.**
- It lands in the Agent 365 registry with its Entra Agent ID — governed from the start.
- Tenant-wide visibility controls decide who can discover and use it.
- This is where “Build” meets “Govern”: the agent we made becomes a governed, shipped product.

THE FULL LOOP

Built in Foundry → grounded with IQ → published & governed in Agent 365. One continuous story.

Agent identity, blueprint, and the Autopilot identity

Three things people conflate — and only the last one is a real Microsoft 365 user.

SERVICE IDENTITY

Foundry Agent Identity

What you see in Foundry today

- Auto-generated by Foundry
- Authenticates the Hosted agent to Azure AI Agent Service
- Runs tools, MCP and Work IQ in technical mode

Not an M365 user — no mailbox, Teams, calendar or Graph.

THE RECIPE

Agent Identity Blueprint

How the identity is defined

- The template Foundry uses to mint agent identities
- Sets attributes, grantable permissions and lifecycle
- Governed by policy and Purview

Still not an M365 user.

REAL M365 IDENTITY

Autopilot Identity

The digital coworker — created at publish to Agent 365

- Full M365 user: mailbox, calendar, Teams, OneDrive, SharePoint, Graph
- Invited to meetings, emailed, @mentioned in Teams
- RBAC + Conditional Access; managed like an employee

This is the identity Work IQ uses.

THE BOTTOM LINE

Your Foundry identity is a technical blueprint — not a Copilot, not a digital coworker, not an Autopilot. The real Microsoft 365 identity is created only when you publish the agent to Agent 365.

Two identities by design — what Work IQ sees

The identity your agent runs with decides what it is allowed to see.

WHAT WORK IQ SEES

✗ **FOUNDRY IDENTITY ONLY**

No meetings. The identity has no calendar, mailbox or Graph permissions.

✓ **PUBLISHED TO AGENT 365**

Work IQ sees the Autopilot's meetings — it is a real M365 user.

WHY MICROSOFT SEPARATES THEM

Same split as a service principal (technical) versus a user identity (functional).

Keeping them apart prevents:

- privilege escalation
- accidental access to M365 data
- Conditional Access bypass
- incomplete audit and compliance gaps

THE MAPPING

Foundry Agent Identity → technical execution, Azure AI Agent Service only. **Autopilot Identity** → Microsoft 365 interaction: Graph, mailbox, calendar, Teams.

Publish to Teams & see it in Agent 365

- From the agent, choose Publish → Microsoft Teams / M365 Copilot.
- Open the Microsoft 365 admin center → Agents: the agent appears in the registry.
- Show its owner, source (Foundry) and Entra Agent ID.
- Note the governance controls: visibility, block/quarantine, policies.

ON SCREEN

The agent we built, now published and already governed — no extra registration.

Build → Ground → Govern

BUILD

Foundry Agent Service: a managed runtime, open tools, memory, voice, identity.

GROUND

The four IQs: Work, Foundry, Web and Fabric IQ ground agents in your reality.

GOVERN

Agent 365: publish, secure, observe and manage the whole fleet.

One agent → a grounded agent → a governed fleet. That's the whole platform in one line.

TAKEAWAYS

What to remember — by role

For leadership (Group A)

- Agents are an operating-model shift, not a feature
- Microsoft extends your trusted stack to agents
- Governance & ROI are built in with Agent 365
- The plan is current, not changed — just renamed

For technical teams (Group B)

- A managed runtime removes the production plumbing
- Open standards (MCP/A2A/OpenAPI) prevent lock-in
- Grounding with IQ makes answers correct & cited
- Identity (Entra Agent ID) is first-class from day one

Day 2 — four hands-on labs (for the technical team)

Lab 1 · Create a prompt agent

Build an agent in the Foundry portal: model, instructions, playground.

Lab 2 · Add an MCP tool

Connect a tool via the Toolbox / MCP; invoke it from the agent.

Lab 3 · Call via Responses API

Use each agent's Coding panel: run it from code (Python / REST).

Lab 4 · Publish to Agent 365

Publish and invoke with the agent's own identity — without OBO.

45–60 minutes each · step-by-step scripts · follow along on your own machine

THANK YOU

Questions & discussion

Today: Build → Ground → Govern. Tomorrow: we build it, together.

